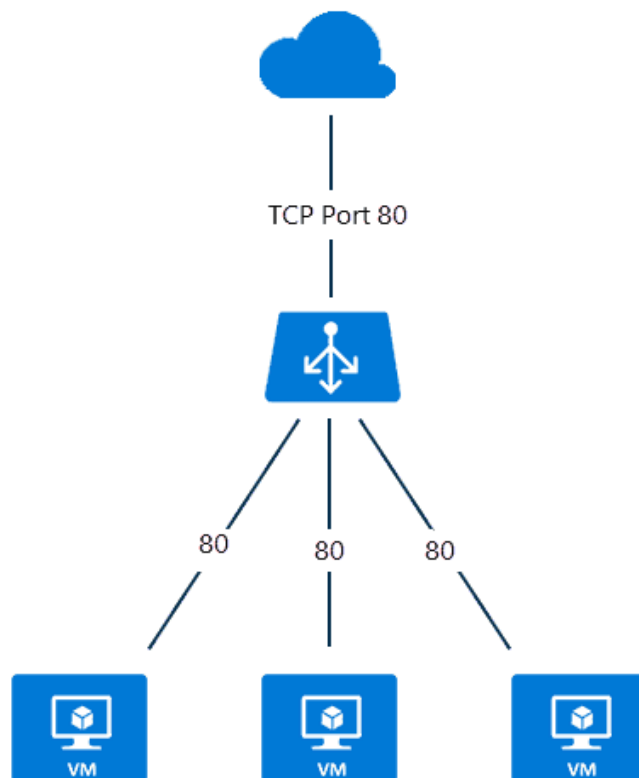


Azure Load Balancer

Azure load balancing is used to distribute traffic loads to backend virtual machines, By defining your own load balancing rules you can use a load balancer in a more flexible way.

What Is Azure Load Balancer?

Azure load balancer allows you to distribute traffic to your backend virtual machines. An Azure load balancer provides high availability for your application. The Azure load balancer is fully managed service itself.



Why Choose Azure Load Balancer?

With Standard Load Balancer, you can scale your applications and create highly available services. Load balancer supports both inbound and outbound scenarios. A load balancer provides low latency and high throughput and scales up to millions of flows for all TCP and UDP applications. Some of the key scenarios that you can accomplish using Standard Load Balancer include:

- Load balance internal and external traffic to Azure virtual machines.
- Increase availability by distributing resources within and across zones.
- Use health probes to monitor load-balanced resources.


- Employ port forwarding to access virtual machines in a virtual network by public IP address and port.
- Standard Load Balancer provides multi-dimensional metrics through Azure Monitor. These metrics can be filtered, grouped, and broken out for a given dimension. They provide current and historic insights into the performance and health of your service. Resource Health is also supported.
- Load balance services on multiple ports, multiple IP addresses, or both.

Load Balancer Types

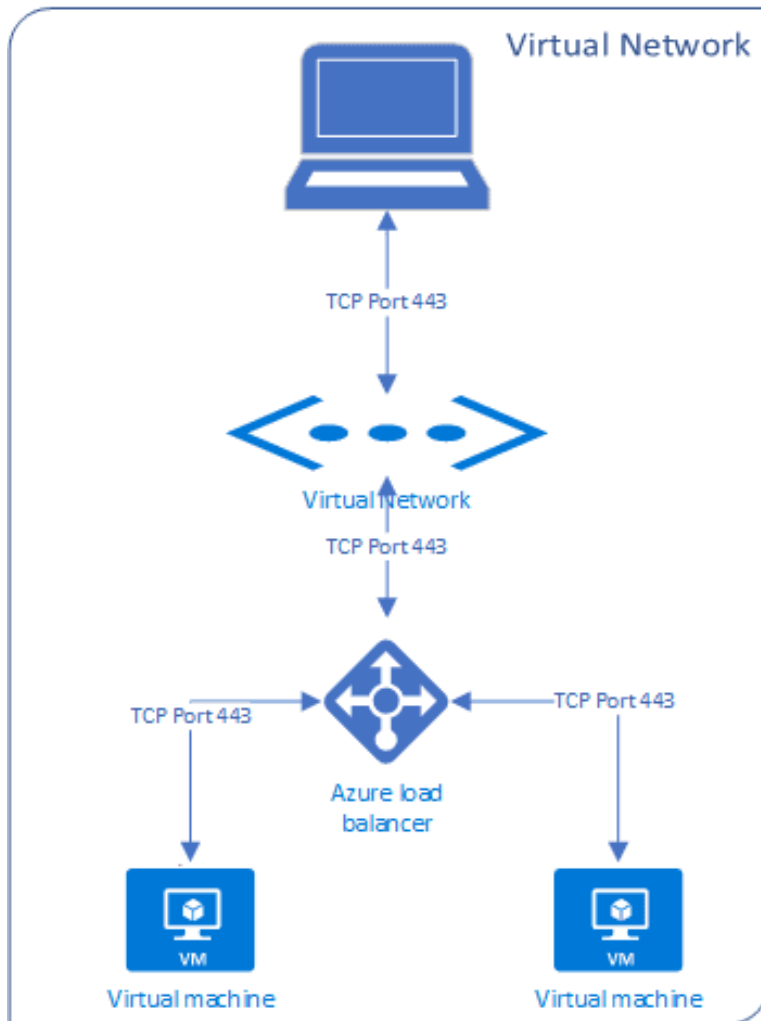
In Azure, you can create two types of the **load balancer**

- Public load balancer
- Internal/ private load balancer

Public Load Balancer: A **public load balancer** can be used to load balance **internet traffic** to virtual machines. It can provide **outbound connections** for virtual machines (VMs) inside your virtual network.



Internal/ Private Load Balancer: An **internal (or private) load balancer** is used to balance traffic from **within a virtual network**.



Along with load balancer, there are two pricing tiers available **Basic** and **Standard**


Basic: Basic tier load balancer provides basic features and restricted to some limits like for backend pool size it is restricted to only 300 instances, it's restricted to a single availability set and the Basic tier only supports multiple frontends for inbound traffic.

Standard: Standard tier load balancer is generally available and offers higher-scale and new features. The Standard tier load balancer is a paid-for feature using a complex set of consumption-based charges and the Basic tier continues to be free. The standard tier can scale out to 1000 instances and the Standard tier can span any virtual machine in a single virtual network, including blends of scale sets, availability sets, and machines.

To know more features of **Basic** and **Standard** tier check the below table

	Standard Load Balancer	Basic Load Balancer
Backend pool size	Supports up to 1000 instances.	Supports up to 300 instances.
Backend pool endpoints	Any virtual machines or virtual machine scale sets in a single virtual network.	Virtual machines in a single availability set or virtual machine scale set.
Health probes	TCP, HTTP, HTTPS	TCP, HTTP
Health probe down behavior	TCP connections stay alive on an instance probe down and on all probes down.	TCP connections stay alive on an instance probe down. All TCP connections end when all probes are down.
Availability Zones	Zone-redundant and zonal frontends for inbound and outbound traffic.	Not available
Diagnostics	Azure Monitor multi-dimensional metrics	Azure Monitor logs
HA Ports	Available for Internal Load Balancer	Not available
Secure by default	Closed to inbound flows unless allowed by a network security group. Internal traffic from the virtual network to the internal load balancer is allowed.	Open by default. Network security group optional.
Outbound Rules	Declarative outbound NAT configuration	Not available
TCP Reset on Idle	Available on any rule	Not available
Multiple front ends	Inbound and outbound	Inbound only
Management Operations	Most operations < 30 seconds	60-90+ seconds typical
SLA	99.99%	Not available

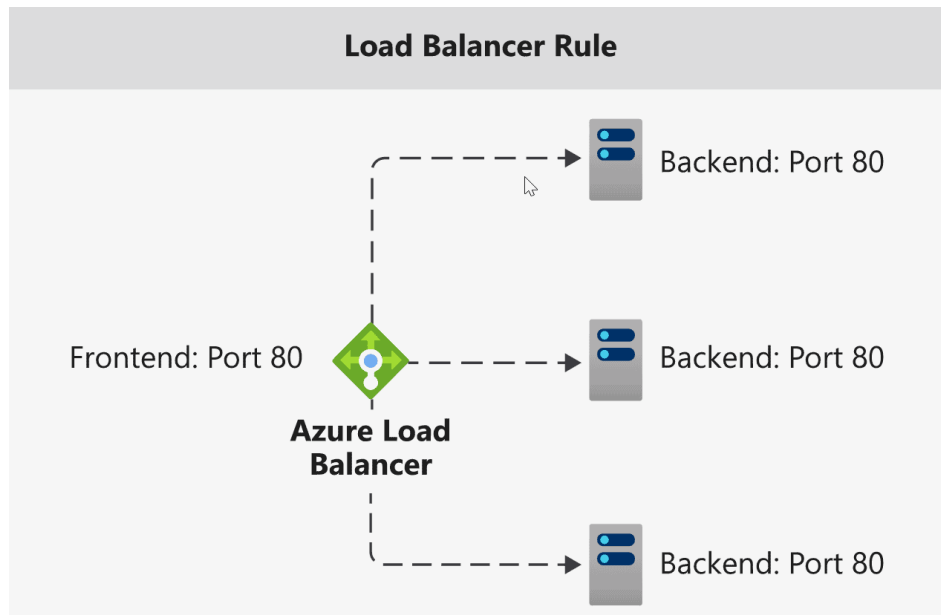
Features Of Azure Load Balancer



- **Load Balancing:** Azure load balancer uses a 5-tuple hash which contains source IP, source port, destination IP, destination port, and protocol.
- **Outbound connection:** All the outbound flows from a private IP address inside our virtual network to public IP addresses on the Internet can be translated to a frontend IP of the load balancer.
- **Automatic reconfiguration:** Load balancer is able to reconfigure itself when it scales up or down instances on the basis of conditions. So, if more virtual machines are added into the backend pool, automatically load balancer will reconfigure.
- **Application agnostic and transparent:** It doesn't directly interact with TCP or UDP protocols. We can route the traffic based on URL or multi-site hosting
- **Health probes:** When any failed virtual machines in a load balancer are recognized by health probe in the backend pool then it stop routing the traffic to that particular failed virtual machine. It can configure a health probe to determine the health of the instances in the backend pool.
- **Port forwarding:** The load balancer supports port forwarding ability if we have a pool of web servers, and we don't want to attach public IP address for the every web server in that pool.

Azure Load Balancing Rules

1. Load balancing rules define how to route traffic when it arrives on the load balancer
2. These rules can use to redirect traffic into the backend pool
3. You can also enable session persistence- Here's client IP's can be directed to the same backend virtual machines



Azure Load Balancer Components

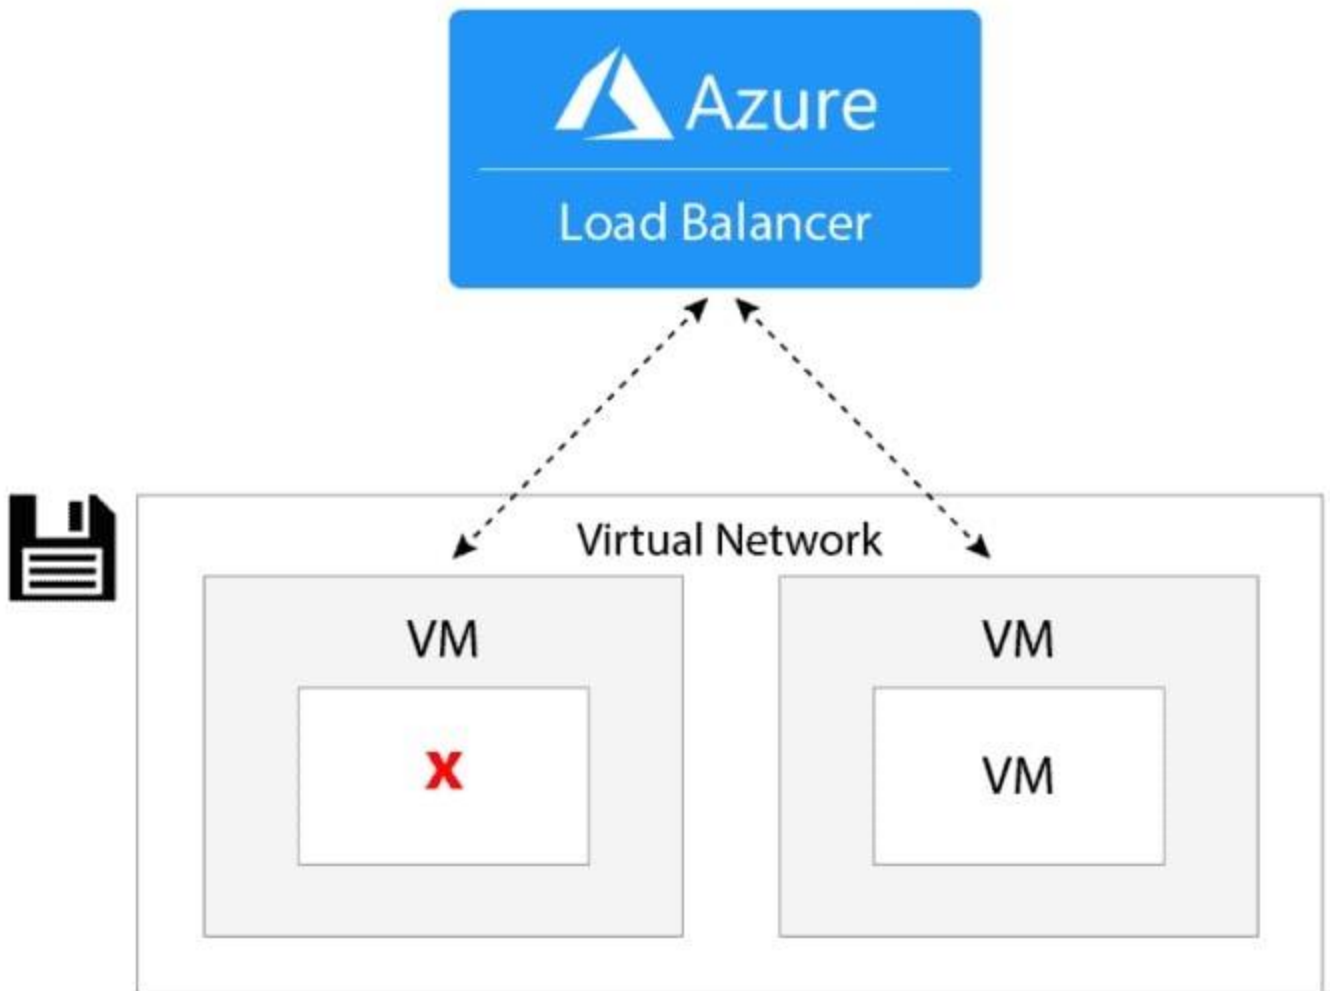
1) Front-end IP configuration: This is the IP address of your Azure load balancer. It is the point of contact for clients. These IP addresses can be:-

- **Public IP Address**
- **Private IP Address**

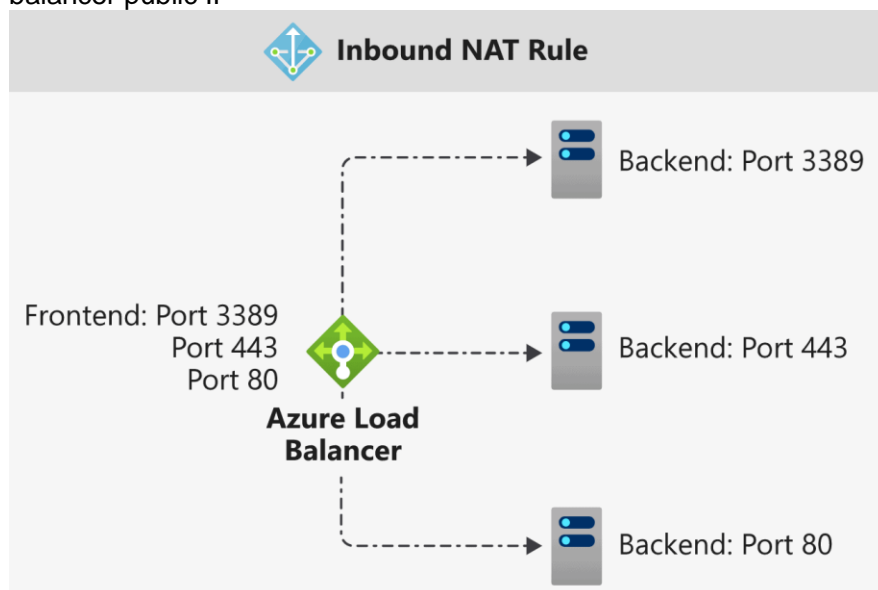
The nature of the IP address explains the sort of load balancer created. Selection of Private IP address gives an internal load balancer. whereas, selection of Public IP address provides public load balancer.

2) Back-end address pool: These are the pool of virtual machines to which the traffic will generally pass. By adding more instances to backend pool it is easy to achieve scale-cost and satisfy high volumes of incoming traffic.

3) Health Probes: A health probe is used to check the health of the backend virtual machine is good or not



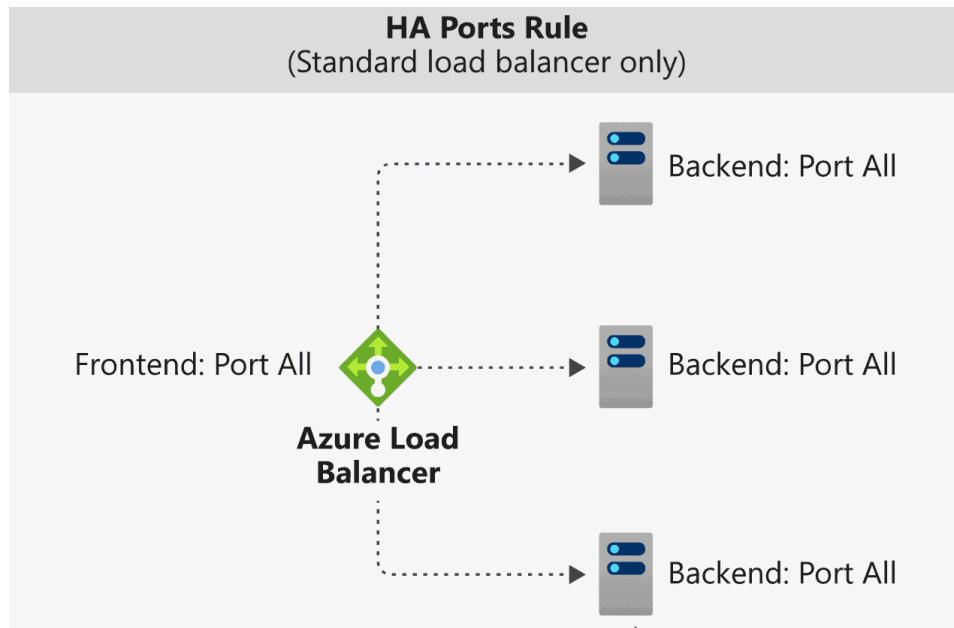
4) Inbound & Outbound NAT rules: By flowing through the front end IP and distributes to the backend IP, NAT rules define the inbound traffic. Outbound rules will transmit VM private IP to load balancer public IP



5) High Availability Ports: The load-balancing decision is made per flow. This action is based on the following five-tuple connection:

1.
 1. source IP address
 2. source port
 3. destination IP address
 4. destination port
 5. protocol

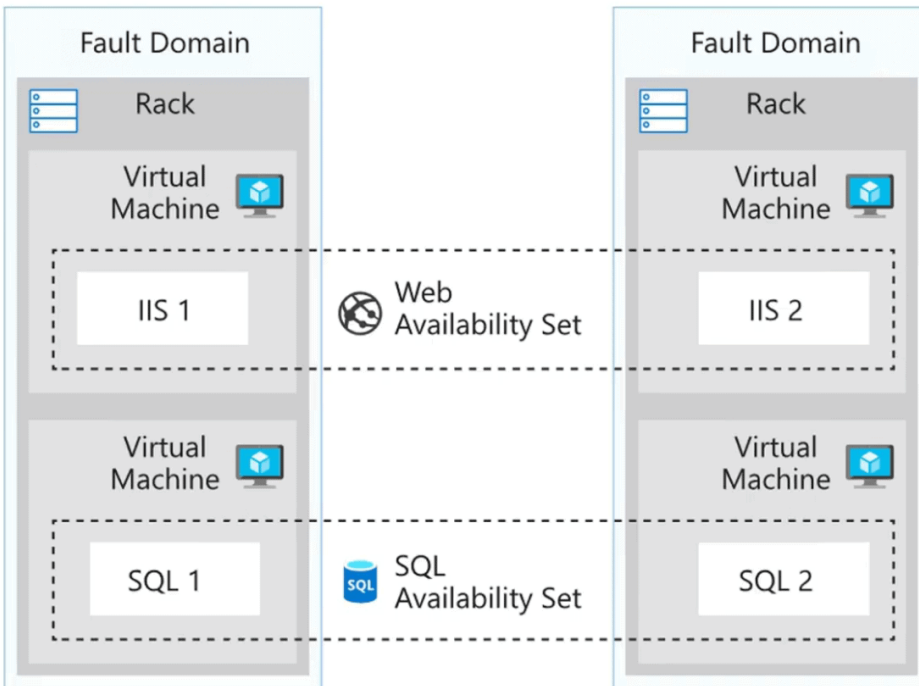
The HA ports load-balancing rules assist you with critical scenarios, like high availability and scale for network virtual appliances (NVAs) inside virtual networks.



Select a Load Balancer Solution

Basic Load Balancers allow:

- Port Forwarding
- Automatic Reconfiguration
- Health Probes
- Outbound Connections through source network address translation (SNAT)
- Diagnostics through Azure Log Analytics for public-facing load balancers.
- It can be used in **availability sets**, not availability zones.



Standard Load Balancers allow:

- HTTPS health probes
- Diagnostics through Azure Monitor, for multidimensional metrics
- High availability (HA) ports
- Outbound rules
- A guaranteed SLA (99.99% for two or more Virtual Machines in different zones)
- It can be used in **availability zones**

Azure Region

